

The `rredis` Package

Bryan W. Lewis
blewis@illposed.net

October 20, 2011

1 Introduction

The `rredis` package provides a native R interface to Redis. Redis is an in memory key/value database with many innovative features written by Salvatore Sanfilippo. It supports data persistence, networked client/server operation, command pipelining, structured value types, server replication, data expiration, clustering, multicast-like publish/subscribe, and it's very fast.

The following simple example illustrates a typical use of the `rredis` package:

```
> library('rredis')
> redisConnect()
> redisSet('x', rnorm(5))
[1] TRUE
> redisGet('x')
[1] 0.808448325 0.341482747 -0.728739322 -0.105507214 -0.002349064
```

The key name “x” is associated with the R vector produced by `rnorm(5)` and stored in Redis. Note that the R object associated with “x” is available to other `rredis` clients, and indeed to any Redis client that can deserialize R objects. Neither the Redis server nor the `rredis` clients need reside on the machine on which the result was generated. Depending on the Redis server settings, “x” can be persistent—that is the value and its association with “x” will persist even if the Redis server is terminated and re-started.

Values in Redis are classified by type. Value types are perhaps the most distinguishing feature of Redis.

- The canonical *string* type holds general-purpose objects, for example any serializable R object, text, or arbitrary binary data.

- The *list* type represents lists of Redis *string* objects, ordered by insertion order. Data can be accessed from lists with stack-like PUSH and POP operations, or by directly indexing ranges of elements. Importantly, redis lists support atomic blocking and asynchronous operation.
- Redis *sets* are unordered collections of unique Redis *strings* equipped with typical set operations like unions and intersections. Uniqueness is enforced by Redis at insertion-time. Redis also supports operations on ordered sets, with set commands prefixed by “Z.”
- Redis *hashes* are collections of Redis *strings* indexed by a hashed set of secondary keys.

Expiration intervals or absolute expiration times may be set on any Redis value. The Redis server can handle lots of small transactions with aplomb, easily exceeding 50,000 transactions/second even on very limited hardware¹. Although Redis is primarily an in-memory database, it uses a virtual memory scheme to support large objects and databases larger than available RAM.

2 Supported Platforms

The Redis server is written in ANSI C and supported on most POSIX systems including GNU/Linux, Solaris, *BSD, and Mac OS X. The server is not officially supported on Windows systems at the time of this writing (March, 2010).

The `rredis` package for R is supported on all supported R platforms, including Microsoft Windows, and can connect to a Redis server running on a supported platform.

Redis clients are available for lots of languages other than R, including Java, C, C#, Ruby, Python, PHP, Tcl, Perl, Erlang, Clojure, Javascript, Scala, and more...

2.1 Obtaining and Installing the Redis server

Redis is an open-source project available from <http://redis.io>, with source code available from Github at <http://github.com/antirez/redis>. Redis is also available as an installable package for most modern GNU/Linux operating systems.

The Redis server is completely configured by the file `redis.conf`. In order to run the Redis server as a background process, edit this file and change the line:

```
daemonize no
```

to:

```
daemonize yes
```

¹Redis easily exceeds 100,000 transactions/second on typical high-end workstations

Most default Redis configuration files are set to disconnect connected clients after an inactivity time out interval. It's possible to disable that behavior in the `redis.conf` file with:

```
timeout 0
```

You may wish to peruse the rest of the configuration file and experiment with the other server settings as well. Finally, start up the Redis server with

```
redis-server ./redis.conf
```

Note that some packaged versions of Redis on GNU/Linux set the Redis server to start on boot as a service.

3 The `rredis` Package by Example

We explore operation of many of the Redis features available to R through a few examples. Seek out the `rredis` package documentation and the excellent Redis Wiki referenced therein for additional help and examples.

3.1 Basic Operation and Redis Strings

Redis *strings* represent the canonical value type. They are used to store any R object that can be serialized to a bit-stream. Most R objects are serializable. Notable exceptions include objects with open connections and external reference pointers.

We assume from now on that the `rredis` package is loaded in the running R session using either `require('rredis')`

or

```
library('rredis')
```

prior to running any example.

Open a connection to a Redis server with `redisConnect`. By default, `redisConnect()` attempts to connect to a Redis server locally on a default port (6379). Explicitly specify a host and/or port to connect to a server running on a computer different from the computer on which the R session is running, for example,

```
redisConnect(host='illposed.net', port=5555)
```

to connect to a Redis server running on host 'illposed.net' at port 5555.

Once connected we can easily store and retrieve values in the Redis database with `redisSet` and `redisGet`:

```
> x <- rnorm(5)
> print(x)
[1] -0.3297596 1.0417431 -1.3216719 -0.8186305 -0.2705817
> redisSet('x',x)
[1] TRUE
> y <- redisGet('x')
> print(y)
[1] -0.3297596 1.0417431 -1.3216719 -0.8186305 -0.2705817
> all.equal(x,y)
[1] TRUE
> redisGet('z')
NULL
```

Note that one must explicitly specify a key name (“x” in the above example) and that Redis key names need not correspond to R variable names.

The SET/GET operations are atomic—that is, multiple SET and or GET operations are guaranteed not to simultaneously occur. And `redisGet` always returns immediately, even if a value is not available in which case it returns NULL (see the example).

The true power of Redis becomes apparent when we share values across multiple clients. For example, start up a new R session and try:

```
> library('rredis')
> redisConnect()
> y <- redisGet('x')
> print(y)
[1] -0.3297596 1.0417431 -1.3216719 -0.8186305 -0.2705817
```

The default behavior of Redis is to make the database persistent, so the value associated with “x” in the above examples will last until it is overwritten or explicitly removed, even if the Redis server is re-started. One may immediately purge Redis of all key/value pairs with the (dangerous) `redisFlushAll` command.

Redis supports multiple distinct key workspaces, indexed by number. Access may be switched between workspaces with the `redisSelect` function as illustrated below. We also use `redisKeys` to list all key names in the current workspace.

```
> redisKeys()
[[1]]
[1] "x"

> redisSelect(1)
[1] "OK"
> redisKeys()
```

`NULL`

```
redisSelect(0)
> redisKeys()
[[1]]
[1] "x"
```

The number of available workspaces is user-configurable in the `redis.conf` file (the default is 16). Note also that index values in Redis begin with 0.

One may easily store and retrieve multiple objects in one operation with `redisMSet` and `redisMGet`. The example also illustrates how values may be expired (in this case, after one second) with `redisExpire`.

```
> redisMSet(list(x=pi,y=runif(5),z=sqrt(2)))
[1] TRUE
> redisMGet(c('x','y','z'))
$x
[1] 3.141593
$y
[1] 0.85396951 0.80191589 0.21750311 0.02535608 0.11929247
$z
[1] 1.414214

> redisExpire('z',1)
[1] TRUE
> Sys.sleep(1)
> redisGet('z')
NULL
```

3.2 Sharing Data with Clients other than R

Redis provides a particularly convenient system for sharing data between diverse applications. We illustrate cross-application communication with simple examples using R and the `redis-cli` command-line program that is included with the Redis server.

Store a sample value in the Redis database with the `redis-cli` program from the command line as follows:

```
redis-cli set shell "Greetings, R client!"
OK
```

Now, leaving the terminal window open, from an R session, try:

```
> redisGet('shell')
[1] "Greetings, R client!\n"
```

And, voilà, R and shell communicate text through Redis.

The reverse direction requires more scrutiny. From the R session, run:

```
> redisSet('R', 'Greetings, shell client!')
```

And now, switch over to the shell client and run:

```
./redis-cli get R
<<Partially decipherable garbage>>
```

This example produces undesirable results because the default behavior of the R `redisSet` command is to store data as R objects, which the shell client cannot decipher. Instead, we must encode the R object (in this case, a character string) in a format that shell can understand:

```
> redisSet('R', charToRaw('Greetings, shell client!'))
[1] TRUE
```

And now, switch over to the shell client and run:

```
./redis-cli get R
Greetings, shell client!
```

It can be tricky to share arbitrary R objects with other languages, but raw character strings usually provide a reasonable, if sometimes inefficient, common tongue.

The `RAW=TRUE` option may be set on most package functions that receive data, for example `redisGet`. Use the `RAW` option to leave the message data as is (otherwise the functions try to deserialize it to a standard R object). The `RAW` format is useful for binary exchange of data with programs other than R.

3.3 Redis Lists

Redis list value types provide us with a remarkably powerful and rich set of operations. Redis lists may be used to set up data queues and they may be accessed either synchronously or asynchronously.

We walk through basic Redis list operation in the first example below. The example shows how `redisLPush` pushes values onto a list from the left, and `redisRPush` pushes values from the right.

```
> redisLPush('a', 1)
[1] 1
> redisLPush('a', 2)
[1] 2
```

```
> redisLPush('a',3)
[1] 3
> redisLRange('a',0,2)
[[1]]
[1] 3
[[2]]
[1] 2
[[3]]
[1] 1

> redisLPop('a')
[1] 3
> redisLRange('a',0,-1)
[[1]]
[1] 2
[[2]]
[1] 1

> redisRPush('a','A')
[1] 3
> redisRPush('a','B')
[1] 4
> redisLRange('a',0,-1)
[[1]]
[1] 2
[[2]]
[1] 1
[[3]]
[1] "A"
[[4]]
[1] "B"

> redisRPop('a')
[1] "B"
```

Like the `redisGet` function, `redisLPop` and `redisRPop` always return immediately, even when no value is available in which case they return `NULL`. Redis includes a blocking variant of the list “Pop” commands that is illustrated in the next example.

```
> redisBLPop('b',timeout=1)
NULL

> redisLPush('b',runif(5))
```

```
[1] 1
> redisBLPop('b', timeout=1)
$b
[1] 0.3423658 0.4188430 0.2494071 0.9960606 0.5643137
```

In the first case above, the NULL value is returned after a one-second timeout because no value was immediately available in the list. Once populated with data, the second attempt consumes the list value and returns immediately.

We can also block on multiple lists, returning when data is available on at least one of the lists:

```
> redisFlushAll()
[1] "OK"
> redisLPush('b', 5)
[1] 1
> redisBLPop(c('a', 'b', 'c'))
$b
[1] 5
```

Although blocking list operations seem simple, they provide an extraordinarily powerful environment for coordinating events between multiple R (and other client) processes. The following example illustrates a simple event stream in which data is emitted periodically by a shell script, and consumed and processed as events arrive by an R process.

First, open an R window and block on the “a” and “b” lists:

```
> redisFlushAll()
> for (j in 1:5) {
+ x <- redisBLPop(c('a', 'b'))
+ print (x)
+ }
```

Your R session should freeze, waiting for events to process.

Now, open a terminal window and navigate to the directory that contains the `redis-cli` program. Run (the following may all be typed on one line):

```
for x in 1 2 3 4 5;do sleep $x;
  if test $x == "2";
    then ./redis-cli lpush a $x;
    else ./redis-cli lpush b $x;
  fi;
done
```

And now you will see your R session processing the events as they are generated by the shell script:


```
$b  
[1] "1"
```

```
$a  
[1] "2"
```

```
$b  
[1] "3"
```

```
$b  
[1] "4"
```

```
$b  
[1] "5"
```

Now, imagine that events may be processed independently, and that they occur at an extraordinary rate—a rate too fast for R to keep up. The solution in this case is simple, start up another R process and it will handle events as they come in, relieving the first R process of about half the event load. Still not enough, start up another, etc.

Keeping in mind that the R clients can run on different computers, we realize that this simple example can easily lead to a very scalable parallel event processing system that requires very little programming effort!

3.4 Redis Sets

The Redis set value type operates somewhat like Redis lists, but only allowing unique values within a set. Sets also come equipped with the expected set operations, as illustrated in the following example.

```
> redisSAdd('A',runif(2))  
[1] TRUE  
> redisSAdd('A',55)  
[1] TRUE  
> redisSAdd('B',55)  
[1] TRUE  
> redisSAdd('B',rnorm(3))  
[1] TRUE  
> redisSCard('A')  
[1] 2  
> redisSDiff(c('A','B'))  
[[1]]  
[1] 0.5449955 0.7848509
```

```
> redisSInter(c('A','B'))
[[1]]
[1] 55

> redisSUnion(c('A','B'))
[[1]]
[1] 55

[[2]]
[1] 0.5449955 0.7848509

[[3]]
[1] -1.3153612 0.9943198 -0.3725513
```

Redis sets do not include blocking operations.

4 Transactions

Redis supports batch submission of multiple Redis operations. Aggregating operations with transactions can in many cases significantly increase performance. The following description is adapted from the Redis documentation at <http://redis.io>:

The `redisMulti`, `redisExec`, `redisDiscard` and `redisWatch` form the foundation of transactions in Redis. They allow the execution of a group of commands in a single step, with two important guarantees:

1. All the commands in a transaction are serialized and executed sequentially. It can never happen that a request issued by another client is served in the middle of the execution of a Redis transaction. This guarantees that the commands are executed as a single atomic operation.
2. Either all of the commands or none are processed. The `redisExec` command triggers the execution of all the commands in the transaction, so if a client loses the connection to the server in the context of a transaction before calling the `redisMulti` command none of the operations are performed, instead if the `redisExec` command is called, all the operations are performed. When using the append-only file Redis makes sure to use a single `write(2)` syscall to write the transaction on disk. However if the Redis server crashes or is killed by the system administrator in some hard way it is possible that only a partial number of operations are registered. Redis will detect this condition at restart, and will exit with an error. Using the `redis-check-aof` tool it is possible to fix the append only file that will remove the partial transaction so that the server can start again.

Queued Redis commands may be discarded with the `redisDiscard` function. Upon successful execution of a transaction, the results from all of the queued commands are returned as a list.

The `redisWatch` function provides a check and set style conditional transaction. Use `redisWatch` to monitor any number of Redis keys. If any watched key values change prior to calling `redisExec` the entire queued sequence will be discarded. Conditioning transactions with `redisWatch` is quite useful in multi-client asynchronous settings.

The following extremely basic example illustrates transactions conditional on no change in value corresponding to the “z” key:

```
> redisWatch('z')
[1] "OK"
> redisMulti()
[1] "OK"
> redisSet('x',runif(3))
[1] "QUEUED"
> redisSet('y',pi)
[1] "QUEUED"
> redisGet('x')
[1] "QUEUED"
> redisExec()
[[1]]
[1] "OK"

[[2]]
[1] "OK"

[[3]]
[1] 0.7620601 0.5982853 0.8274721
```

5 Publish/Subscribe

The publish/subscribe functions let Redis clients reliably multicast (publish) messages over “channels” that any client may subscribe to. Channels are identified by name (character string). Use the `redisSubscribe` function to subscribe to one or more channels. Use the `redisPublish` function to transmit messages over a channel. Once subscribed, channels must be monitored for incoming messages using the `redisGetResponse` function, usually in an event loop. Beware that the `redisGetResponse` function indefinitely blocks for an incoming message on subscribed channels.

Here is a simple example:

```
> redisSubscribe('channel1')
```

```
# The loop will receive three messages from 'channel1':
> for(j in 1:3) print(redisGetResponse())
# A typical message might look like:
[[1]]
[1] "message"

[[2]]
[1] "channel1"

[[3]]
[1] "message3"

# Finally, unsubscribe to the channel:
> redisUnsubscribe('channel1')
```

Note that the only Redis functions that may be used in between the `redisSubscribe` and `redisUnsubscribe` functions are `redisGetResponse`, `redisSubscribe`, and `redisMonitorChannels` functions.